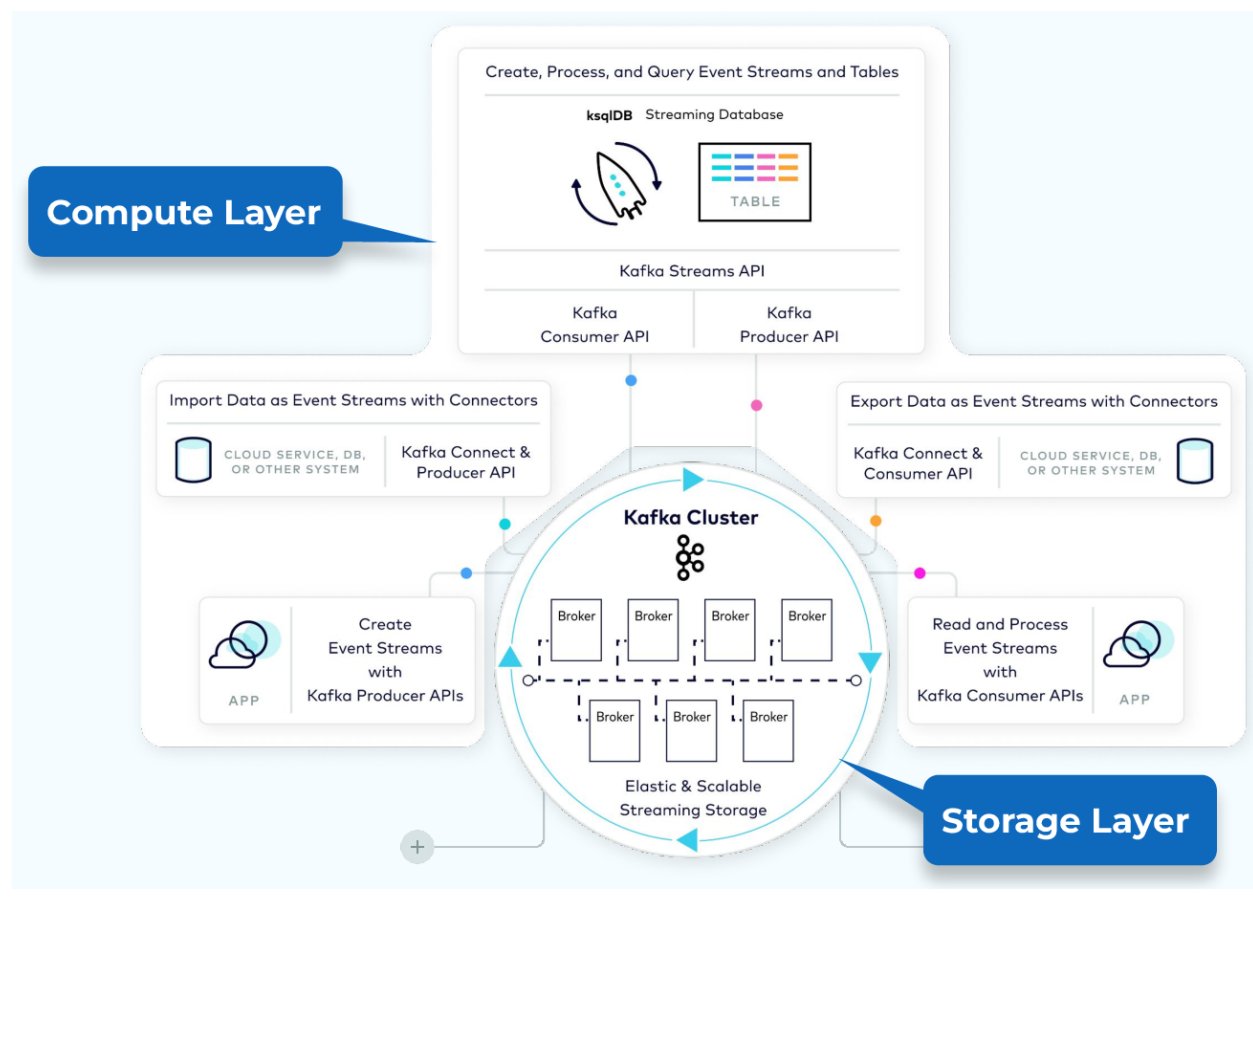


Fundamentals

Overview of Kafka Architecture



- Kafka is a data streaming system that allows developers to react to new events as they occur in real time
 - Kafka architecture consists of a storage layer and a compute layer.
 - The storage layer is designed to store data efficiently and is a distributed system such that if your storage needs grow over time you can easily scale out the system to accommodate the growth
 - The compute layer consists of four core components—the producer, consumer, streams, and connector APIs, which allow Kafka to scale applications across distributed systems.
-

Producer and Consumer APIs

The foundation of Kafka's powerful application layer is two primitive APIs for accessing the storage.

- the producer API for writing events
 - the consumer API for reading them.
-

Kafka Connect

Kafka Connect, which is built on top of the producer and consumer APIs, provides a simple way to integrate data across Kafka and external systems.

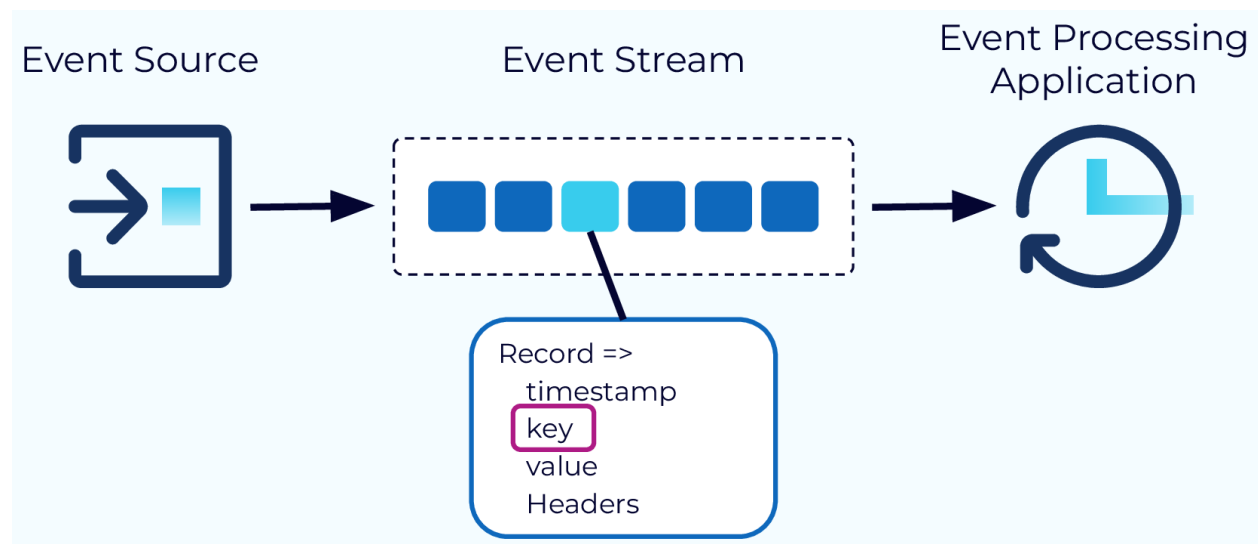
- Source connectors bring data from external systems and produce it to Kafka topics.
 - Sink connectors take data from Kafka topics and write it to external systems.
-

Kafka Streams

For processing events as they arrive, we have Kafka Streams, a Java library that is built on top of the producer and consumer APIs.

Kafka Streams allows you to perform real-time stream processing, powerful transformations, and aggregations of event data.

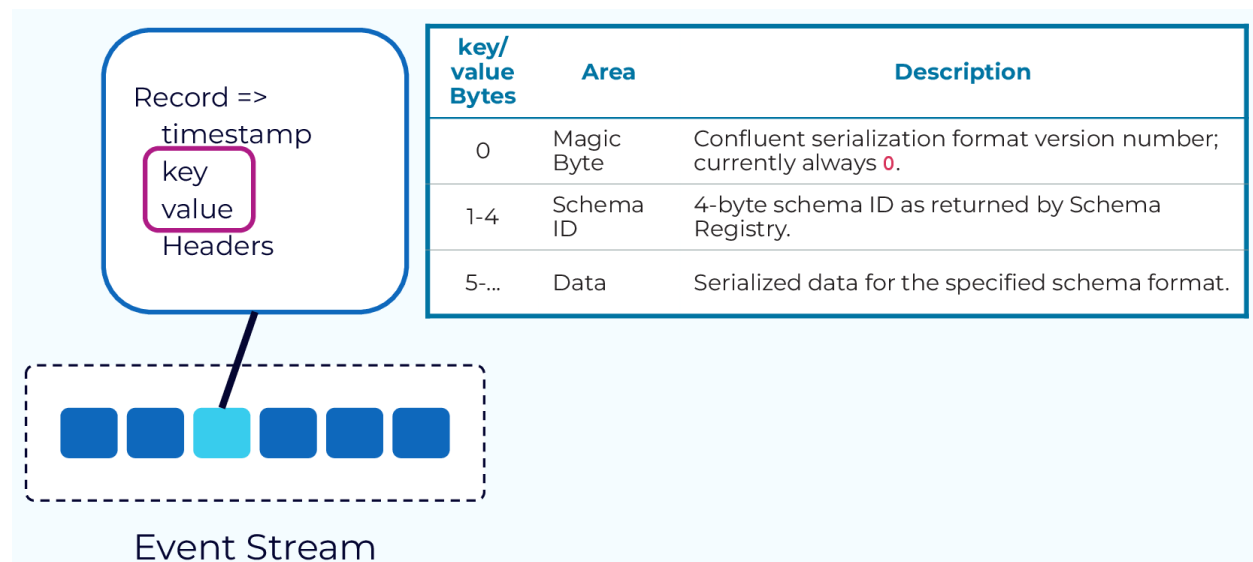
What Is an Event in Stream Processing?



- An event is a record of something that happened that also provides information about what happened.
- Examples of events are customer orders, payments, clicks on a website, or sensor readings.
- An event shouldn't be too large.
- A 10GB video is not a good event. A reference to the location of that video in an object store is.

- An event record consists of a timestamp, a key, a value, and optional headers.
- The event payload is usually stored in the value.
- The key is also optional, but very helpful for event ordering, colocating events across topics, and key-based storage or compaction.

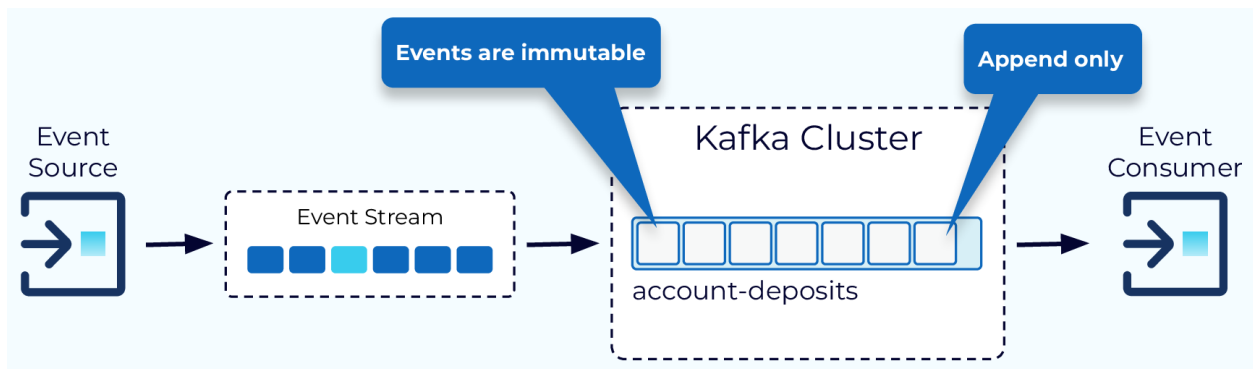
Record Schema



- In Kafka, the key and value are stored as byte arrays which means that clients can work with any type of data that can be serialized to bytes.
- A popular format among Kafka users is Avro, which is also supported by Confluent Schema Registry.
- When integrated with Schema Registry, the first byte of an event will be a magic byte which signifies that this event is using a schema in the Schema Registry.

- The next four bytes make up the schema ID that can be used to retrieve the schema from the registry, and the rest of the bytes contain the event itself.
- Schema Registry also supports Protobuf and JSON schema formats.

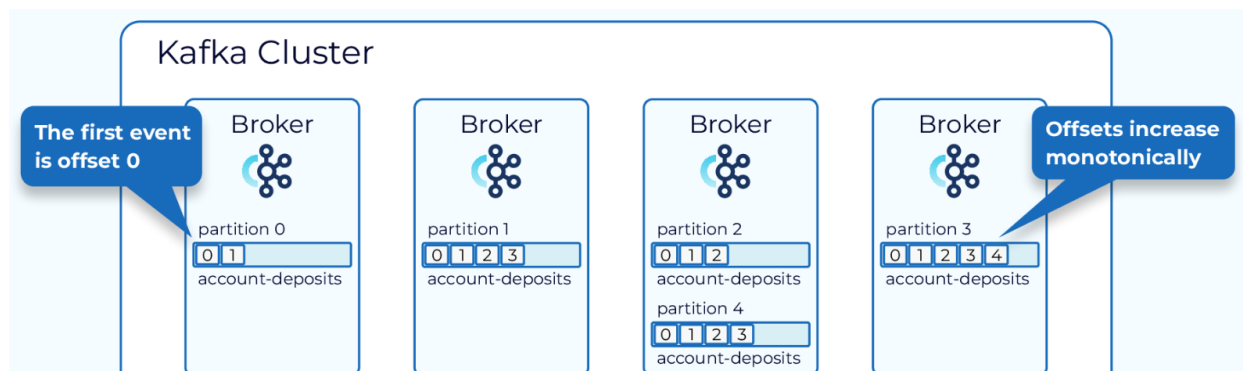
Kafka Topics



- Named container for similar events
 - System contains lots of topics
 - Can duplicate data between topics
- Durable **logs** of events
 - Append only
 - Can only seek by offset, not indexed
- Events are **immutable**

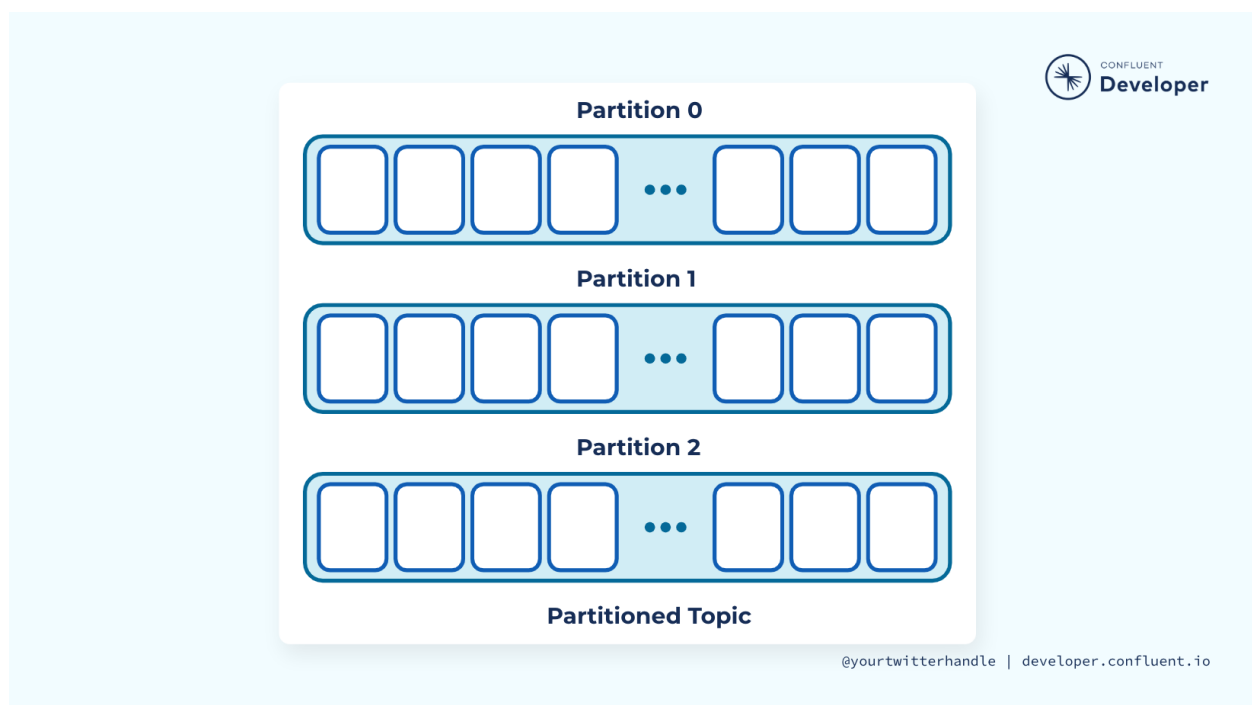
- A key concept in Kafka is the topic.
- Topics are append-only, immutable logs of events.
- Typically, events of the same type, or events that are in some way related, would go into the same topic.
- Kafka producers write events to topics and Kafka consumers read from topics.

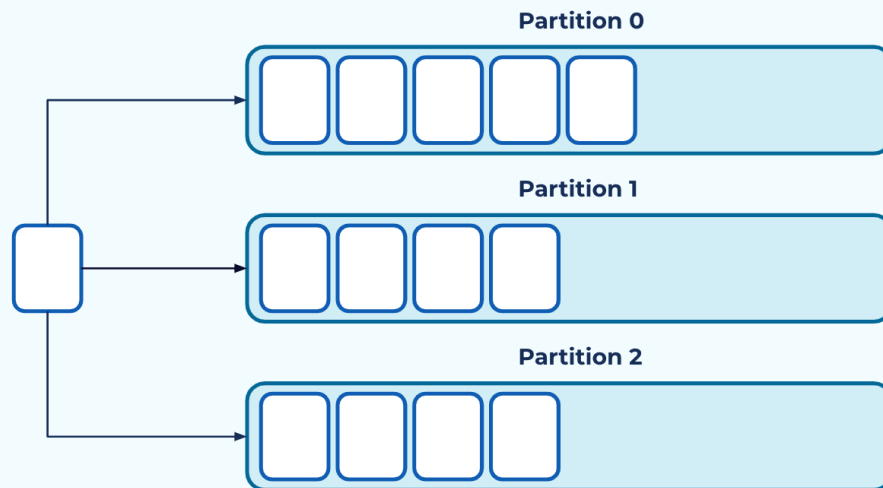
Kafka Topic Partitions



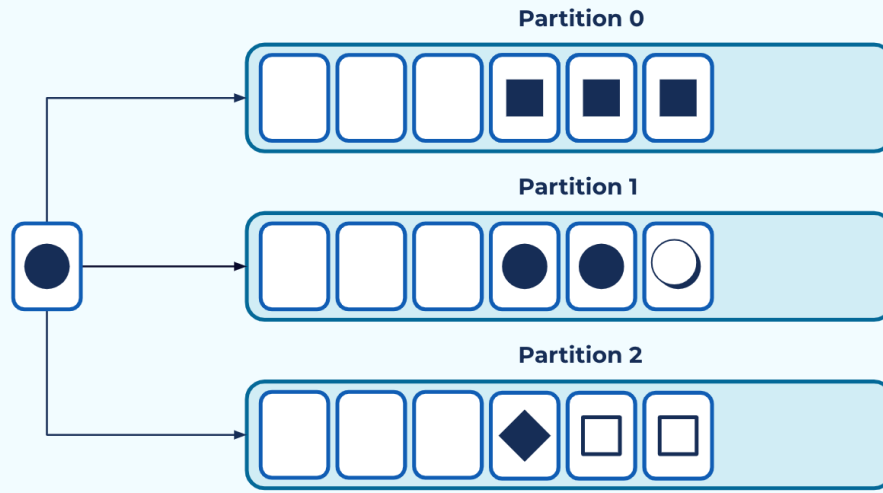
- In order to distribute the storage and processing of events in a topic, Kafka uses the concept of partitions.
- A topic is made up of one or more partitions and these partitions can reside on different nodes in the Kafka cluster.

- The partition is the main unit of parallelism.
 - Events can be produced to a topic in parallel by writing to multiple partitions at the same time.
 - Likewise, consumers can spread their workload by individual consumer instances reading from different partitions. If we only used one partition, we could only effectively use one consumer instance.
-
- Within the partition, each event is given a unique identifier called an offset.
 - The offset for a given partition will continue to grow as events are added, and offsets are never reused.
 - The offset has many uses in Kafka, among them are consumers keeping track of which events have been processed.






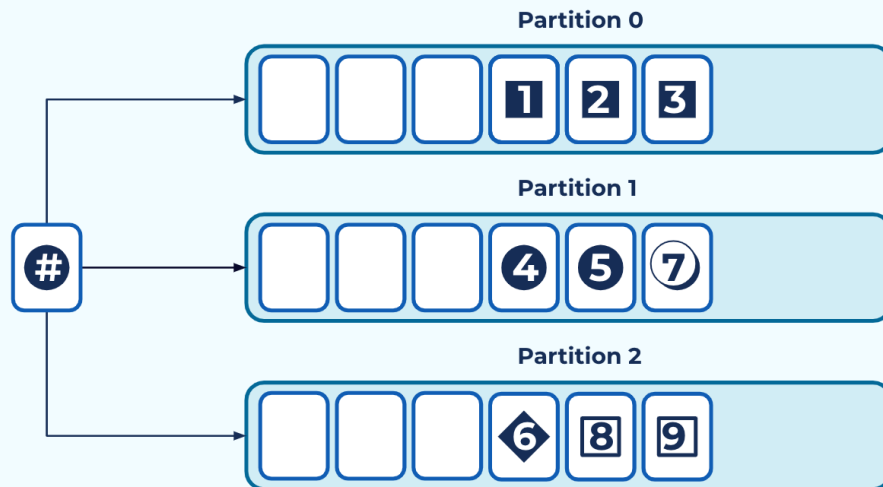
@yourtwitterhandle | developer.confluent.io



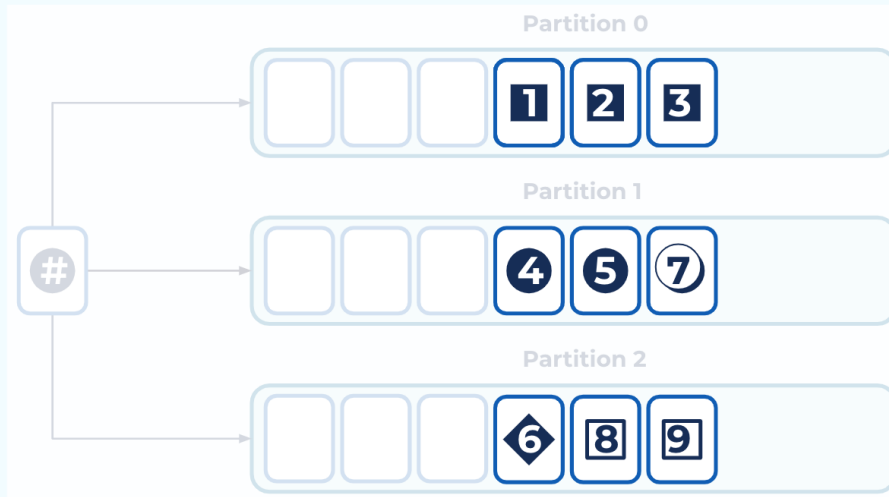
@yourtwitterhandle | developer.confluent.io



@yourtwitterhandle | developer.confluent.io



@yourtwitterhandle | developer.confluent.io



@yourtwitterhandle | developer.confluent.io

Producer is to decide which partition to send the messages to.

Depending on various configuration and parameters, the producer decides the destination partition.

Lets check different scenario:-

1. **No key specified** :- producer will randomly decide partition and would try to balance the total number of messages on all partitions.
2. **Key Specified** :- the producer uses **Consistent Hashing** to map the key to a partition. consistent is a hashing mechanism where for the same key same hash is generated always, and it minimizes the redistribution of keys on a re-hashing scenario
3. **Partition Specified** :- You can hardcode the destination partition as well.

4. **Custom Partitioning logic:-** We can write some rules depending on which the partition can be decided.
-